

overview

today's spend, throughput, cache and rate-limit headroom across your local sessions.

filters | all providers | all models

period

today

block

7d


14d

30d

90d

180d

all



live · 01HXKQ3F · leibniz-platform/main

sonnet 4.5

ctx


48%

94k / 200k

\$0.34

2m 41s

→



0

refusals

3



max_tokens cuts

2

retry loops

1

errors



rate limits

5h cap 38%

1.9M / 5.0M resets 2h 14m

7d cap 87%

34.8M / 40.0M resets 4d 6h

● daemon · running

14ms

session	preview	project	model	flags	turns	cost	last
● 01HXKQ3F	add idempotency key to /sessions endpoint, wire through store layer	leibniz-platform	● sonnet 4.5		42	\$2.41	now
● 01HXJZ7R	port supplier-sync cron to async, add backoff, test against prod replica	leibniz-platform	● opus 4	retry	88	\$5.12	12m ago
● 01HXJYWQ	refactor pricing-rule engine; was getting opus to think out loud about rules	leibniz-platform	● opus 4	max_tok	31	\$3.04	42m ago
● 01HXJYAB	codex: scan repo for unused exports, write a prune patch	autoresearch	● gpt-5 · codex		18	\$0.62	1h ago
● 01HXJX9P	investigate flaky test_session_replay; suspect race on close()	leibniz-platform	● sonnet 4.5	err	24	\$0.94	2h ago
● 01HXJWKE	polish onboarding copy, generate hero variants	acme-web	● haiku 4.5		14	\$0.18	3h ago
● 01HXJVT2	upgrade ruff config and migrate deprecated checks	internal-tools	● sonnet 4.5		9	\$0.21	5h ago

L

leibniz

v0.4.2

workspace

dashboard

sessions86

settings

add idempotency key to /sessions endpoint

sonnet4.5

live · 2m41s

health

max_tokens · 1

retry · 2

refusals · 0

errors · 0

trace

context window

transcript

- workspace
- dashboard
- sessions 86
- settings

rate limits

5h cap38%1.9M / 5.0Mresets 2h 14m

7d cap87%34.8M / 40.0Mresets 4d 6h

● daemon · running 14ms

sessions / 01hvkq3f · leibniz-platform/main

add idempotency key to /sessions endpoint

● sonnet 4.5


● live · 2m 41s

open jsonl


health max_tokens · 1 retry · 2 refusals · 0 errors · 0

trace context window transcript

a / context evolution · % used over turns · amber marks compaction peak 99% · current 48%



b / composition · 94k of 200k · 48% used



- agents 12k
- skills 18k
- mcp 9k
- memory 22k
- tools 14k
- conv. 19k

c / categorized · what's actually loaded · token cost per item

custom agents · 1	12k	skills · 3	18k
router-architect	12.0k	testing-conventions	7.4k
		api-style-guide	6.2k
		migration-runner	4.4k

mcp servers · 2	9k	memory · 4 files	22k
github · 8 tools	5.8k	~/CLAUDE.md	8.1k
postgres · 4 tools	3.2k	leibniz/CLAUDE.md	6.4k
		PRINCIPLES.md	4.8k
		ARCHITECTURE.md	2.7k

system tools · 14k

read	2.1k	edit	1.9k	grep	1.4k	bash	2.7k
glob	0.9k	task	1.4k	webfetch	1.1k	todowrite	2.5k